Title: Reasoning Approaches in AI Models

Research Intern Take-Home Assignment

Author: Mahima Kumari
Date: 27/04/2025

1 . Introduction

In order to ensure safety, trust, and ethical alignment with human values, it is crucial to comprehend the reasoning mechanisms of AI systems as they become more and more integrated into industries like healthcare, banking, education, and governance. AI that is transparent and explicable encourages user accountability and confidence, while early bias detection and mitigation avoid biased results and advance equity. AI that possesses strong reasoning skills can also manage complicated and unclear situations, which enhances the quality of decision-making in crucial applications like financial forecasts and diagnostics. In addition, ethical frameworks and transparency that integrate AI in search with human principles support responsible deployment and the well-being of society. In along with improving scientific abilities, this strategy guarantees AI functions as a reliable, human-centered partner, reducing the possibility of inadvertent harm and promoting the responsible adoption of game-changing technology.

2. Individual Reasoning

2.1 Deductive reasoning

Deductive reasoning is the process of drawing specific conclusions from general, most likely true premises or principles. This top-down approach makes sense in some situations by using generally accepted fundamental principles. Deductive reasoning is renowned for its truth-preserving properties, which guarantee that if the initial premises are true, the conclusion is also true. Common signs of deductive thinking include syllogisms and other structured logical structures where two or more premises logically lead to a conclusion (e.g., "All humans are mortal; Socrates is a human; therefore, Socrates is mortal"). Modus ponens and modus tollens are two more formal structures that are essential to mathematical logic and computer science. Formal logic engines and rule-based expert systems in artificial intelligence (AI) are based on deductive reasoning, which is the sequential application of predefined facts and rules to infer new information. This approach is highly advantageous for tasks requiring rigorous correctness and formal verification, such as mathematical proofs, automated theorem proving, legal reasoning, and safety-critical decision-making systems (such those in medical diagnostics or flight control). However, strictly logical AI models may find it challenging to handle open-ended or ambiguous situations, where necessary premises may not be stated explicitly, or where knowledge is unclear or incomplete. In these circumstances, other forms of reasoning, such as abductive or inductive reasoning, might be more appropriate.

Since deductive reasoning is predicated on logical necessity the idea that the conclusion is the only possible outcome given the premises, it is very reliable in precisely defined domains. Its shortcomings, however, highlight how important it is to integrate deductive reasoning methods with other reasoning approaches in complex, useful AI systems. Deductive reasoning fails primarily due to mistakes or a lack of knowledge. The deductions won't be allowed if the rules or premises are off. Another weakness of the system is its incapacity to manage circumstances that aren't specifically addressed by the standards. The selection and development of the underlying assumptions, which represent the viewpoints or constraints of the knowledge engineers, may lead to biases. In fields with set standards and a need for accuracy, Because of its logical structure and intrinsic certainty, deductive reasoning is secure and reliable for applications like formal verification and regulatory compliance. However, it might not be as useful in situations that are more open-ended or unclear due to its rigidity and requirement for complete data.

2.2 Inductive Reasoning

Using a bottom-up approach, inductive reasoning draws generalizations from specific facts or instances. Artificial intelligence systems can use this method to infer more general principles from patterns in data. In contrast to deductive thinking, induction draws conclusions that are probabilistic, meaning that they are likely true but not necessarily so given the available information. As a result, inductive reasoning embraces ambiguity and draws lessons from empirical data rather than solely depending on accepted assumptions. Inductive reasoning, in which models are educated on data patterns and extrapolate from existing examples to predict future events, is heavily utilized in machine learning algorithms. One important use of inductive reasoning is supervised learning, which allows AI systems to learn to predict outcomes using labeled training data. Many real world applications Use this approach, such as spam filtering. With this method, algorithms search through enormous volumes of tagged emails for traits that could indicate spam. Similar to this, inductive reasoning is applied to picture classification problems, where models learn visual characteristics from known ones to accurately classify new, unlabeled photos. In Natural Language Processing (NLP) tasks such as text classification and sentiment analysis, which likewise heavily rely on inductive reasoning, large-scale corpora of labeled text are processed to discover important patterns for language understanding. Personalized experiences on streaming and e-commerce platforms are also made possible by recommendation systems, which utilize inductive reasoning to predict user preferences based on historical activity.Inductive reasoning is demonstrated when multiple white swans are observed and the conclusion that all swans are white is reached. Consider an artificial intelligence (AI) image recognition system that has been trained on thousands of labeled images of dogs and cats. By inductively learning general characteristics that distinguish cats from dogs through the examination of factors such as ear shape, tail length, and face anatomy across samples, the model is able to classify new, unseen photos. Similar to this, a spam filter looks for trends in a significant number of emails that have been reported as spam, such as commonly used keywords or questionable sender addresses

Furthermore, inductively trained models may suddenly fail due to distributional shift, which occurs when the properties of fresh data diverge dramatically from those of the training data. The probabilistic character of inductive reasoning and its dependence on data have important ramifications for the security and reliability of AI. Although flexibility and strong generalization are made possible by inductive learning, the concerns of bias, injustice, and unreliability demand careful design decisions. For safety-critical fields including healthcare, banking, and autonomous systems, rigorous validation, bias prevention, and ongoing monitoring are crucial. Making sure that training data is balanced, varied, and representative of the real-world deployment environment is essential to developing reliable inductive reasoning-based AI systems.

2.3 Abductive Reasoning

Abductive reasoning, a kind of logical inference, begins with an incomplete collection of observations and applies logical inference to ascertain the most plausible explanation for them. The best explanation is inferred from observed evidence, even if there is no assurance of certainty. "Inference to the best explanation," or abductive reasoning, is the process of developing hypotheses to explain observed events and then evaluating them in light of the available data to ascertain which is the most logical and likely. Abductive reasoning is often used in artificial intelligence diagnostic systems. Finding the most likely cause of a problem from a collection of potentially inadequate data is the aim of medical diagnosis tools and fault detection systems.(AI) systems use abduction to decipher the intended meanings of unclear phrases or statements in a particular context. In order to identify trends suggestive of fraudulent activity in vast amounts of transactional data, the banking industry uses abductive reasoning in AI systems. The street is wet" is an observation to think about. The most likely explanation, based on abductive reasoning, is that it rained last night, particularly if combined with a weather forecast that indicates rain. In a medical setting, an AI system may assume that a heart condition is the most likely cause of a patient's severe chest pain, shortness of breath, and dizziness, especially if the patient has a family history of heart disease and high cholesterol, according to their electronic health records. In a similar vein, when a self-driving car spots an unexpected impediment, it may employ abductive reasoning to evaluate several potential causes, such as a pedestrian, debris, or a sensor malfunction. It will then test each theory in light of new sensor data and past trends Inadequate or inaccurate prior knowledge can result in failure modes in abductive reasoning, which can cause the choice of a tenable but ultimately incorrect explanation. If a less obvious explanation does not fit well with what the system already knows, it may ignore it. Inaccurate conclusions can also result from biases in the assessment of hypotheses, which may be based on preconceived notions or the organization of the information base. Because abductive reasoning is inherently unpredictable, AI safety and reliability must be carefully considered. The lack of certainty and the possibility of contradicting theories necessitate thorough validation and human supervision, especially in safety-critical fields like healthcare, even though its capacity to produce plausible explanations is useful in many applications, including diagnostics. Building reliable abductive reasoning systems requires that the AI's knowledge base be complete, accurate, and bias-free.

2.4 Analogical Reasoning

Drawing comparisons between various circumstances or ideas in order to draw conclusions or find solutions is known as analogical reasoning. The way it works is by identifying parallels between a new, unknown scenario and a previously understood, familiar one, and then applying the knowledge from the familiar situation to the new one. This procedure enables AI systems to discover underlying structural correspondences, transfer knowledge across disciplines, and produce innovative solutions. Analogical reasoning is very helpful in case-based reasoning systems in artificial intelligence. By accessing comparable previous examples and modifying their solutions to fit the current situation, these systems are able to tackle new problems. By identifying commonalities between seemingly unrelated ideas, analogical reasoning can help improve machine learning algorithms by allowing them to generalize from sparse data. It is useful in the representation and transfer of knowledge, enabling AI to apply knowledge acquired in one field to another with similar characteristics. AI can benefit from analogical reasoning in a number of ways. It makes it possible to draw conclusions from sparse data by facilitating generalization from restricted data. It encourages efficiency and creativity by making it easier for information and approaches to problem-solving to be transferred between other fields. Analogical thinking can produce original and creative solutions that may not be obvious using standard reasoning techniques by establishing links between concepts that don't seem to be connected. By connecting new and complicated circumstances to ones that are already known, it facilitates comprehension and makes knowledge more useful and accessible. Analogical reasoning failure types include focusing on superficial parallels while ignoring the core distinctions between the source and target domains. It is possible to draw incorrect conclusions by extrapolating from a small number of similar examples. The choice of analogies that fit the AI's prior "knowledge" or the biases in its training data are two examples of bias. Analogical thinking plays a complicated role in the safety and reliability of AI. Its capacity to apply current knowledge to new issues can help with situational adaptation, but the possibility of faulty analogies casts doubt on the validity of the conclusions, particularly in situations when safety is a top priority. Building confidence in AI systems' analogical reasoning abilities requires that they be able to recognize profound, pertinent parallels while avoiding deceptive analogies.

2.5 Causal Reasoning

Finding cause-and-effect links between variables or occurrences is known as causal reasoning. Causal reasoning seeks to ascertain if one event directly effects another, in contrast to conventional statistical techniques that mostly concentrate on correlations. AI systems can reason about interventions ("What happens if we change X?") and counterfactuals ("Would Y have occurred if X had been different?") thanks to this type of reasoning, which goes beyond simple prediction and helps them comprehend the fundamental mechanics that govern outcomes.

Causal reasoning is a technique used in artificial intelligence to create systems that understand underlying mechanisms and make decisions based on them rather than just obvious patterns. AI may use causal models in the healthcare industry to ascertain whether a certain treatment directly enhances patient outcomes. To investigate the impact of shifting product rankings on user behavior, recommendation systems might employ causal reasoning. Since causal models can be used to find real causes of outcomes rather than phony correlations, addressing bias and enhancing the resilience of AI systems are also important applications. These linkages are formally formalized using methods such as Pearl's "do-calculus," Bayesian networks, and structural equation models. Causal AI is also used in root cause analysis, causal effect estimation, causal fairness, and algorithmic recourse (offering optimal solutions). Inaccurate causal assumptions can result in failure modes in causal reasoning, which can produce faulty conclusions. Unmeasured confounders can skew causal estimates since they affect both the cause and the effect. Biased judgments about causal linkages can also result from biases in the data used to learn those correlations. Models that ignore intricate or indirect causal paths may be erroneous or incomplete.In terms of AI safety and reliability, causal reasoning is extremely pertinent since it focuses on comprehending real relationships and facilitating reasoning about interventions. Building trust and guaranteeing safety in high-stakes applications requires more dependable, equitable, and transparent systems, which causal AI can produce by going beyond prediction to explanation and comprehension.

2.6 Chain-of-Thought (CoT) Reasoning

Large language models (LLMs) can improve their reasoning skills by being encouraged to deconstruct difficult problems into a sequence of intermediate, logical steps through the use of Chain-of-Thought (CoT) prompting. The model is asked to openly demonstrate its reasoning process rather than giving a definitive response, thereby simulating the methodical approach that people frequently take while completing complex tasks. By guiding the LLM through intermediate thinking phases, this technique improves its ability to solve challenging tasks like math problems, symbolic manipulation, and commonsense reasoning. Applications for CoT prompting can be found in a variety of prompting techniques for LLMs, such as faithful CoT (which gaurntee that the final response follows the reasoning chain), contrastive CoT (which displays both correct and incorrect reasoning), few-shot CoT (where the prompt includes examples of reasoning steps), zero-shot CoT (where the model is simply instructed to "think step by step"), and automatic CoT (which generates reasoning examples automatically). Tasks like creating summaries, evaluating compliance documents, and combining data from many sources have all been accomplished with success. CoT can be demonstrated with a math word problem such as "John has 10 apples." After giving away four, he gets five more. "How many apples is he holding?" The model's reasoning would be guided by a CoT prompt: "John starts with 10 apples." Given that he gives away 4, 10 - 4 = 6. Five more apples are then given to him, making 6 + 5 = 11. The final response is 11". The model clearly displays the preliminary computations that result in the ultimate response.

"Let's think step by step" is a simple prompt that may be used in zero-shot CoT to encourage the model to come up with its own reasoning after posing a complex question. Illogical or factually inaccurate reasoning processes are examples of CoT failure mechanisms, especially in smaller models. To support a response, models may produce reasoning that seems reasonable but is ultimately incorrect. The reasoning process may be impacted by biases in the training data. Additionally, models may learn to produce reasoning that maximizes positive feedback, even if it isn't totally accurate, a phenomenon known as reward hacking. CoT prompting has important ramifications for AI safety and reliability because it aims to make LLMs' reasoning process more transparent and reasonable. By offering an insight into the model's "thought" process, it helps foster user confidence and make it easier to see possible biases or mistakes. To guarantee its dependability, particularly in safety-critical applications, more study and improvement is necessary, as evidenced by the worries about the reasoning's faithfulness and its reliance on model size.

2.7 Tree-of-Thoughts (ToT) Reasoning

Tree-of-Thoughts (ToT) prompting is a sophisticated method that allows large language models (LLMs) to investigate and assess several diverse lines of reasoning at once, improving their ability to solve problems. ToT uses a tree structure instead of the linear Chain-of-Thought method, where each node represents a partial answer or an intermediate thought in the reasoning process, and branches represent many possible operators or future steps. This methodology mimics a more human-like approach to difficult problem-solving by enabling the model to retrace its steps when a given course of action appears unlikely to result in a workable solution and to investigate alternate approaches as necessary. idea deconstruction (dividing the problem into smaller steps), idea creation (producing various potential solutions or next steps), and state evaluation (evaluating the potential of each partial solution) are some of the essential elements of ToT prompting. To travel the tree of thoughts, it can use a variety of search algorithms, such as Depth-First Search (DFS) or Breadth-First Search (BFS). In tasks that call for preparation and foresight, such numerical reasoning, creative writing, and puzzles, ToT has been demonstrated to perform better than alternative prompting techniques. ToT prompting has the important benefit of helping LLMs solve problems by allowing them to investigate several lines of reasoning at once, which is comparable to how the brain works. It can look ahead and retrace, which makes it appropriate for situations requiring sophisticated decision-making. Higher success rates and more cohesive outcomes on mentally taxing tasks may emerge from it, according to research. One of the failure modes in ToT is becoming trapped in branches of the reasoning tree that aren't useful, which results in computation waste. Potentially accurate pathways may be prematurely pruned due to biases in the evaluation function used to evaluate the mental states. In areas with limited resources, its practical utility may be limited by its high computing cost. Despite its resource intensity, the ability of ToT to enhance problem-solving by exploring multiple reasoning paths has positive implications for AI safety and trustworthiness. By encouraging a more deliberate and comprehensive exploration of potential solutions, it can potentially lead to more robust and reliable outcomes.

2.8 Graph-based Reasoning

Graph-based reasoning in AI involves representing information as a network of interconnected entities (nodes or vertices) and the relationships between them (edges). This approach emphasizes the connections and context within data, allowing AI systems to move beyond processing isolated data points and perform logical inferences based on the graph structure. These inferences can include pathfinding (identifying sequences of relationships), pattern matching (discovering recurring structures), subgraph analysis (examining local relationships), and rule-based reasoning (applying predefined rules to the graph). Knowledge graphs are an important application of graph-based reasoning. Through methods like Graph Retrieval-Augmented Generation (GraphRAG), these structured representations of knowledge are being incorporated with big language models more frequently in order to provide AI reasoning a factual foundation and improve its correctness and dependability. In order to generate knowledge maps, identify multidisciplinary connections, and propose new lines of inquiry, graph-based reasoning is frequently employed in the analysis of scientific literature. The ability to reason about networked data is essential in fields such as clinical decision support, risk assessment, financial forecasting, and AI-assisted decision-making. Think of a knowledge graph that shows supply chain information. Suppliers, items, and locations could be represented by nodes, while relationships like "supplies" or "located in" could be represented by edges. An AI system may use graph-based reasoning to determine the quickest route between a supplier and a client for a particular product or assess how a disruption at one point in the supply chain would affect the chain as a whole. An artificial intelligence (AI) model may be used in scientific research to examine a knowledge graph of biological materials and their characteristics, finding surprising connections between them or proposing new material designs based on what is already known. A knowledge graph can be used to answer questions by navigating the connections between the items the question mentions.

3. Comparative Analysis of AI Reasoning Approaches

Data requirements and common uses. Expert systems employ deductive reasoning, which uses broad rules to arrive at particular conclusions. Machine learning and prediction are powered by inductive reasoning, which extracts general patterns from specific data. Abductive reasoning is helpful in diagnostics because it creates tenable explanations from insufficient data. Analogical thinking uses analogous situations from the past to answer new challenges. In order to facilitate decision-making in the face of uncertainty, causal reasoning establishes cause-and-effect linkages. By decomposing thinking into steps or examining several possible solutions, the chain-of-thoughts and tree-of-thoughts techniques assist language models in resolving complicated issues. Knowledge graphs and data analysis are supported by graph-based reasoning, which makes use of organized links between items. Every strategy has distinct advantages and works well for certain AI tasks.

Table

Reasoning Approach	Information Processing	Data Requirements	Data Requirements	Key Applications
Deductive	General to Specific	General principles, known premises	Certain	Expert systems, rule-based systems, formal logic
Inductive	Specific to General	Specific instances, observations, data	Probable	Machine learning, pattern recognition, prediction
Abductive	Observations to Plausible Explanation	Incomplete observations, evidence	Probable	Diagnostics, medical diagnosis, fault detection
Analogical	Comparing Similar Situations	Past experiences, similar scenarios	Probable	Case-based reasoning, knowledge transfer, problem-solving
Causal	Identifying Cause and Effect	Data capturing correlations and context	Variable	Decision-making under uncertainty, bias reduction
Chain-of-Thought	Generating Intermediate Reasoning Steps	Prompt with instructions and examples	Variable	LLMs for complex tasks (math, reasoning)
Tree-of-Thoughts	Exploring Multiple Reasoning Paths in a Tree Structure	Prompt with instructions and problem statement	Variable	LLMs for complex problem-solving, planning
Graph-based	Leveraging Interconnected Entities and Relationships	Structured data, knowledge graphs	Variable	Knowledge graphs, question answering, data analysis

4. Strengths and Weaknesses of Reasoning Approaches in AI

This is a thorough table that examines the advantages and disadvantages of each of the reasoning techniques you listed, with an emphasis on how well they work with AI systems. Important topics included in the analysis include precision, interpretability, effectiveness, resilience, and possible failure modes.

Reasoning Approach	Strengths	Weaknesses
Deductive Reasoning	- High accuracy if premises are true - Logically sound and consistent - Easy to trace and interpret - Supports formal verification	- Not robust to noisy/incomplete data - Inflexible in dynamic environments - Cannot generate new knowledge - Limited real-world use
Inductive Reasoning	- Learns from data patterns - Adapts to new/unseen data - Scalable for big data - Foundation of ML algorithms	- Can overfit/underfit - May generalize incorrectly - Prone to training data bias - Less interpretable
Abductive Reasoning	- Generates plausible hypotheses - Useful in diagnosis and NLP - Works with incomplete info - Flexible reasoning	- No guarantee of correctness - Multiple valid conclusions - Depends on prior/domain knowledge - Hard to verify
Analogical Reasoning	- Enables transfer learning - Aids conceptual generalization - Useful in novel problem-solving	- Weak analogies mislead results - Similarity is hard to define - Limited formalism
Causal Reasoning	- Supports “why” and “what-if” analysis -Enables counterfactual thinking - Important for safe/fair AI	- Causal inference is hard - Requires high-quality or interventional data - Computationally costly
Chain-of-Thought (CoT)	- Encourages logical thinking in steps - Improves LLM reasoning accuracy - More interpretable than direct answers	- Errors propagate step-by-step - Requires multiple outputs for reliability - Can be verbose
Tree-of-Thought (ToT)	- Captures structured relationships - Transparent and queryable - Scales well in knowledge-rich domains	- Needs structured/curated input - Hard to update dynamically - Limited with unstructured data
Knowledge Graphs / Graph-based Reasoning	- Combines logic and learning - More robust to uncertainty - Supports explainable AI	- Complex to implement and integrate - Potential scalability issues - May require domain expertise

5. Conclusion

The foundation of intelligent systems is made up of AI reasoning techniques, each of which has unique advantages and disadvantages for practical uses. Deductive reasoning is essential for fields that need strict accuracy, such formal verification, legal compliance, and rule-based expert systems, because it offers certainty and rigor. Its inflexibility and reliance on comprehensive, precise premises, however, restrict its flexibility in situations that are unclear or open-ended. Machine learning, pattern recognition, and prediction tasks are powered by AI's ability to generalize from specific facts through inductive reasoning. Although its probabilistic nature permits adaptability and flexibility, it also carries the potential of bias and inaccuracy in the event that the training data is faulty or not representative. This makes continuous monitoring and cautious data selection crucial, particularly in settings where safety is a top priority.The foundation of flaw detection and healthcare diagnostic systems is abductive reasoning, which is excellent at producing tenable explanations from partial facts. Its strength is in managing ambiguity, but in situations where the stakes are great, it needs strong validation to prevent incorrect conclusions. By comparing similar situations, analogical reasoning enables AI to transfer knowledge across disciplines, encouraging innovation and effective problem-solving. However, if surface-level resemblances are confused with deeper structural correspondences, it may falter, underscoring the importance of selecting analogies carefullyAI can now recognize actual cause-and-effect linkages by using causal reasoning, which goes beyond correlation. This is essential for interventions in intricate systems like healthcare and finance, as well as for reducing prejudice and making sound decisions. However, clear and thoroughly tested causal models are necessary since causal inference is susceptible to confounding variables and model assumptions. The capacity of huge language models to break down complicated issues, investigate several lines of reasoning, and produce clear, sequential solutions is improved by sophisticated techniques like Chain-of-Thought and Tree-of-Thoughts prompting. These methods enhance problem-solving skills and user confidence, but they need close supervision to guarantee computational efficiency and logical integrity.AI can carry out complex inference and information discovery in fields including supply chain analysis, scientific research, and knowledge graphs thanks to graph-based reasoning, which makes use of interconnected data. The caliber and organization of the underlying data determine how effective it is. In conclusion, not every AI problem can be solved by a single reasoning technique. By combining various reasoning techniques, which are backed by openness, moral principles, and human supervision, AI systems are made to be reliable, strong, and consistent with society's ideals. In order to deploy AI safely, responsibly, and effectively across crucial industries and advance both technological capabilities and human well-being, a multifaceted approach is necessary.There isn't a single AI reasoning technique that can handle all of the problems that arise in practical applications. Deductive, inductive, abductive, analogical, causal, chain-of-thoughts, tree-of-thoughts, and graph-based reasoning all have advantages, but when used alone, they also have drawbacks. For instance, inductive and abductive reasoning are more flexible but can add bias or uncertainty if not well controlled, whereas deductive reasoning guarantees strict correctness in clearly defined cases but struggles with ambiguity However, robust, reliable, and human-aligned AI requires more than just technical integration. In order to foster confidence and facilitate accountability, transparency making AI decision-making processes intelligible to users and stakeholder is crucial.

6. References

1. Russell, S., & Norvig, P. (2021). *Artificial Intelligence: A Modern Approach* (4th ed.). Pearson.
2. Pearl, J. (1988). *Probabilistic Reasoning in Intelligent Systems: Networks of Plausible Inference*. Morgan Kaufmann.
3. Brachman, R. J., & Levesque, H. J. (2004). *Knowledge Representation and Reasoning*. Morgan Kaufmann.
4. Reiter, R. (1980). A logic for default reasoning. *Artificial Intelligence*, 13(1–2), 81–132. https://doi.org/10.1016/0004-3702(80)90014-4
5. Ghallab, M., Nau, D., & Traverso, P. (2004). *Automated Planning: Theory and Practice*. Morgan Kaufmann.
6. McCarthy, J. (1980). Circumscription — A form of non-monotonic reasoning. *Artificial Intelligence*, 13(1–2), 27–39. https://doi.org/10.1016/0004-3702(80)90011-9
7. Poole, D. (1993). Logic programming, abduction and probability. *New Generation Computing*, 11(3–4), 377–400. https://doi.org/10.1007/BF03037261
8. Besold, T. R., Garcez, A. d., Bader, S., Bowman, H., Domingos, P., Hitzler, P., ... & Silver, D. L. (2017). Neural-symbolic learning and reasoning: A survey and interpretation. *arXiv preprint arXiv:1711.03902*. https://arxiv.org/abs/1711.03902
9. Levesque, H. J. (1984). A logic of implicit and explicit belief. *AAAI Conference on Artificial Intelligence*.
10. Koller, D., & Friedman, N. (2009). *Probabilistic Graphical Models: Principles and Techniques*. MIT Press.
11. Davis, E., & Marcus, G. (2015). Commonsense reasoning and commonsense knowledge in artificial intelligence. *Communications of the ACM*, 58(9), 92–103. https://doi.org/10.1145/2701413